

# Reinforcement Learning

Jack Lanchantin

November 30, 2015


# Overview

1 Reinforcement Learning

2 Markov Decision Processes

- **Framing of the problem of learning from interaction to achieve a goal.**
- **Agent:** learner and decision maker
- **Environment:** what the learner interacts with (everything outside the agent)
- Agent selects actions and the environment responds to those actions and presents new situations

# Reinforcement Learning



- At each time step  $t$ , the agent receives the environment **state**  $S_t \in \mathcal{S}$ , and the agent then selects an **action**  $A_t \in \mathcal{A}(S_t)$ 
  - $\mathcal{S}$  is the set of possible states (whatever information is available to the agent).
  - $\mathcal{A}(S_t)$  is set of actions available in state  $S_t$
- One time step later, the agent receives a **reward**,  $R_{t+1} \in \mathcal{R} \subset \mathbb{R}$ , and ends up in a new state  $S_{t+1}$

- At each time step, the agent implements a mapping  $\pi_t$  from states to probabilities of selecting each possible action, where  $\pi_t$  is called a **policy**
  - $\pi_t(a|s)$  = probability that  $A_t = a$  if  $S_t = s$

## Reinforcement Learning Objective

The agent's goal is to maximize the total amount of reward it receives over the long run by changing its policy as a result of its experience

- Let the sequence of rewards after time step  $t$  is  $R_{t+1}, R_{t+2}, R_{t+3}, \dots$ , then we want to maximize the return  $G_t$
- The agent chooses  $A_t$  to maximize the discounted return:

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

where  $\gamma$  is the discount rate and  $0 \leq \gamma \leq 1$

- The closer  $\gamma$  is to 1, the more the agent accounts for future rewards

# Markov Property

- Probability of transitioning to new state  $s$ :

$$p(s', r|s, a) = Pr\{S_{t+1} = s', R_{t+1} = r | S_0, A_0, R_1, \dots, S_{t-1}, A_{t-1}, R_t, S_t, A_t\}$$

- Probability with Markov assumption

$$p(s', r|s, a) = Pr\{S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a\}$$

- Markov assumption allows us to predict the next state and expected rewards from knowledge of only the current state
  - Assume that the current state tells us everything we need to know for future (e.g. current state of checker board)

# Markov Decision Processes

- A R.I. task that satisfies the Markov property is called a Markov Decision Process (MDP)
- Given any state  $s$  and action  $a$ , the probability of each possible pair of next state and reward  $s', r$ , is denoted

$$p(s', r|s, a) = Pr\{S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a\}$$

From this probability representation, we can compute anything else we might need to know about the environment



- Expected rewards for state–action pairs:

$$r(s, a) = \mathbb{E}[R_{t+1} | S_t = s, A_t = a] = \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} p(s', r | s, a)$$

- State-transition probabilities:

$$p(s' | s, a) = \Pr\{S_{t+1} = s' | S_t = s, A_t = a\} = \sum_{r \in \mathcal{R}} p(s', r | s, a)$$

- Expected rewards for state–action–next-state triples:

$$r(s, a, s') = \mathbb{E}[R_{t+1} | S_t = s, A_t = a, S_{t+1} = s'] = \frac{\sum_{r \in \mathcal{R}} r * p(s', r | s, a)}{p(s' | s, a)}$$

Richard S. Sutton and Andrew G. Barto. Reinforcement Learning: An Introduction. MIT Press 2015.