

Reinforcement Learning

Jack Lanchantin


November 30, 2015

- 1 Reinforcement Learning
- 2 Markov Decision Processes
- 3 Value Functions

- 1 Reinforcement Learning
- 2 Markov Decision Processes
- 3 Value Functions

- **Framing of the problem of learning from interaction to achieve a goal.**
- **Agent:** learner and decision maker
- **Environment:** what the learner interacts with (everything outside the agent)
- Agent selects actions and the environment responds to those actions and presents new situations

Reinforcement Learning




- At each time step t , the agent receives the environment **state** $S_t \in \mathcal{S}$, and the agent then selects an **action** $A_t \in \mathcal{A}(S_t)$
 - \mathcal{S} is the set of possible states (whatever information is available to the agent).
 - $\mathcal{A}(S_t)$ is set of actions available in state S_t
- One time step later, the agent receives a **reward**, $R_{t+1} \in \mathcal{R} \subset \mathbb{R}$, and ends up in a new state S_{t+1}

Reinforcement Learning

At each step t ,

- The agent:
 - Receives state S_t
 - Receives scalar reward R_t
 - Executes action A_t
- The environment:
 - Receives action A_t
 - Emits state S_t
 - Emits scalar reward R_t



Reinforcement Learning Objective

- If the sequence of rewards after time step t is $R_{t+1}, R_{t+2}, R_{t+3}, \dots$, then we want to maximize the return G_t
- The agent chooses A_t to maximize the discounted return:

$$G_t = \sum_{k=0}^{T-t-1} \gamma^k R_{t+k+1}$$

where γ is the discount rate and $0 \leq \gamma \leq 1$. The closer γ is to 1, the more the agent accounts for future rewards

- **Select actions to maximise future reward**

Outline

- 1 Reinforcement Learning
- 2 Markov Decision Processes**
- 3 Value Functions

Markov Property

- Probability of transitioning to new state s' with reward r :

$$p(s', r|s, a) = Pr\{S_{t+1} = s', R_{t+1} = r | S_0, A_0, R_1, \dots, S_{t-1}, A_{t-1}, R_t, S_t, A_t\}$$

- Probability with Markov assumption:

$$p(s', r|s, a) = Pr\{S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a\}$$

- Markov assumption allows us to predict the next state and expected rewards from knowledge of only the current state
 - Assume that the current state tells us everything we need to know for future (e.g. current state of checker board)

Markov Decision Processes

- A reinforcement learning task that satisfies the Markov property is called a Markov Decision Process (MDP)
- Given any state s and action a , the probability of each possible pair of next state and reward (s', r) , is denoted

$$p(s', r|s, a) = \Pr\{S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a\}$$

- From this probability representation, we can compute anything else we might need to know about the environment...

- State-transition probabilities:

$$p(s'|s, a) = \Pr\{S_{t+1} = s' | S_t = s, A_t = a\} = \sum_{r \in \mathcal{R}} p(s', r | s, a)$$

- Expected rewards for state-action pairs:

$$r(s, a) = \mathbb{E}[R_{t+1} | S_t = s, A_t = a] = \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} p(s', r | s, a)$$

- Expected rewards for state-action-next-state triples:

$$r(s, a, s') = \mathbb{E}[R_{t+1} | S_t = s, A_t = a, S_{t+1} = s'] = \frac{\sum_{r \in \mathcal{R}} r * p(s', r | s, a)}{p(s' | s, a)}$$

- At each timestep t , a can-collecting robot decides whether it should (1) actively search for a can, (2) remain in place and wait for someone to bring it a can, or (3) go back to charging station to recharge its battery

Outline

- 1 Reinforcement Learning
- 2 Markov Decision Processes
- 3 Value Functions**

- At each time step, the agent implements a mapping π from states to probabilities of selecting each possible action, where π is called a **policy**
 - $\pi_t(a|s) = \text{probability that } A_t = a \text{ if } S_t = s$

Reinforcement Learning Objective

The agent's goal is to maximize the total amount of reward it receives over the long run by changing its policy as a result of its experience

- Almost all RL algorithms involve estimating value functions
- **Value Functions:** functions of states (or state–action pairs) that estimate how good it is for the agent to be in a certain state (or to perform an action in a given state).
 - "How good" is defined in terms of future rewards that can be expected (i.e. expected return)

State-Value Function for policy π

- Recall that a policy π maps a state s and action a to the probability $\pi(a|s)$ of doing a when in s
- The value of a state s under policy π , denoted $v_\pi(s)$, is the **expected return when starting in state s and following π thereafter**:

$$v_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s] = \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s \right]$$

$\mathbb{E}_\pi[\cdot]$ is the expected value of a r.v. given the agent follows policy π

Action-Value Function for policy π

- The value of taking action a in state s under a policy π , denoted $q_\pi(s, a)$, is the **expected return starting from s , taking the action a , and following policy π thereafter**:

$$q_\pi(s) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a] = \mathbb{E}_\pi \left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s, A_t = a \right]$$

- Both v_π and q_π can be estimated from experience using Monte Carlo methods

Computing Value Functions

- Key idea of reinforcement learning is to use value functions to search for a good policy π
- Can use dynamic programming techniques to compute optimal value functions, and thus find optimal policies

$$\begin{aligned}v_{\pi}(s) &= \mathbb{E}_{\pi}[G_t | S_t = s] \\&= \mathbb{E}_{\pi}\left[\sum_{k=0}^{\infty} \gamma^k R_{t+k+1} \mid S_t = s\right] \\&= \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) [r + \gamma v_{\pi}(s')]\end{aligned}$$

Optimal Value Functions

- Solving a reinforcement learning task means finding a policy that achieves a high reward over the long run
- A policy π is better than π' if its expected return is \geq to that of π' for all states
 - I.e. $\pi \geq \pi'$ iff $v_\pi(s) \geq v_{\pi'}(s)$ for all $s \in \mathcal{S}$

Optimal Value Functions

- The *optimal* policies are denoted π_*
- The *optimal state-value functions* are denoted v_* and defined:

$$v_*(s) = \max_{\pi} v_{\pi}(s) \text{ for all } s \in \mathcal{S}$$

- The *optimal action-value functions* are denoted q_* and defined:

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a) \text{ for all } s \in \mathcal{S} \text{ and } a \in \mathcal{A}$$

Richard S. Sutton and Andrew G. Barto. Reinforcement Learning: An Introduction. MIT Press 2015.